



Mudasir Shafiq

DATA & SOFTWARE ENGINEER

CONTACT

- +33 77573 96 00
- chmudasir.fr@gmail.com
- Paris, France
- Portfolio
- LinkedIn
- GitHub

SKILLS

Streaming and Pipelines

Kafka Airflow Spark Hadoop

Cloud

AWS GCP Azure Databricks

Warehousing

BigQuery Snowflake Redshift

Transformation

dbt Data Vault 2.0 SQL

Full Stack

React Django Spring Boot PHP

Languages

Python SQL Java Scala

DevOps

Docker Kubernetes Terraform CI/CD

BI and Analytics

Power BI Tableau Grafana GA4/GTM

ML

scikit-learn pandas numpy

EDUCATION

MSc Data Engineering

Aivancity University, Paris
2023 2024

BS Mathematics

Air University
2016 2020

CERTIFICATIONS

- AWS Cloud Practitioner
- IBM Data Engineer
- Google Data Analytics
- Machine Learning Coursera

LANGUAGES

English Advanced
French A2 / B1

PROFILE

Results-driven Data and Software Engineer with 4+ years of experience building real-time streaming pipelines, cloud data warehouses, and full-stack web applications. MSc Data Engineering, Aivancity University Paris. Expert in Kafka, Airflow, BigQuery, dbt and multi-cloud architectures (AWS / GCP / Azure). Full-stack capability across React, Django, Spring Boot with a proven track record delivering end-to-end solutions for real clients across AgriTech, FinTech, Web3, and e-commerce.

EXPERIENCE

Data Engineer

2023 2024

Ceres-Advisory (Blockchain Web3) Paris

- Architected real-time GA4 to BigQuery streaming pipeline via GTM, reducing reporting latency from hours to minutes across 15+ live KPI dashboards
- Built Power BI and Grafana dashboards for live campaign monitoring, enabling instant decisions with zero manual exports
- Engineered blockchain data models with optimised SQL (partitioning, window functions), cutting query time by 60% on multi-million-row datasets
- Designed serverless AWS architecture (Lambda, Step Functions, S3) for resilient ingestion with zero pipeline failures post-deployment

Data Analyst and Performance Marketing

2021 2023

PriceOye (E-Commerce) Pakistan

- Maintained ROAS between 6.33x and 43x, achieving GMV target of PKR 18M per month through data-driven campaign optimisation
- Grew retention rate from 10% to 15% via AI-powered audience segmentation and lifecycle campaigns
- Engineered automated Power BI pipelines for real-time ROI tracking, eliminating 10+ hours per week of manual reporting

KEY PROJECTS

Livestock Intelligence

Kafka • Airflow • BigQuery • scikit-learn

- Real-time dairy herd pipeline: Kafka ingest into BigQuery with CTE feature engineering, 5 anomaly flags and LAG window functions for herd-relative benchmarking
- 3 production ML models covering health classification, fertility prediction and milk-yield regression

Synthetic Data Engine

Python • pandas • Google Sheets API

- Rule-based KPI generator disaggregating 22 months of history into 626 daily rows via Dirichlet weighting with a 12-check validator
- Idempotent cron-scheduled Sheets upsert with monotonicity, engagement ordering and fee-rate correctness checks

Reddit ETL Pipeline

Airflow • Terraform • AWS • dbt • Redshift • Tableau

- Automated end-to-end ETL on AWS following Data Vault 2.0 with dbt transformations and modelling in Amazon Redshift
- Interactive Tableau dashboards delivering analytics and business reporting to stakeholders

Delivery Route Optimisation

Django 5 • Python • SQLite • OpenStreetMap

- Open-source fleet management for SMBs: role-based access, one-click route optimiser and geocoded live map over OpenStreetMap
- UUID public tracking portal, CSV bulk import and export with status-change email notifications

Digital Agency Platform

React 18 • Spring Boot 3 • PostgreSQL • Docker • CI/CD

- Database-driven SPA with 11 typed REST endpoints; multi-stage Docker builds deployed via GitHub Actions on every push to main
- Zero-deploy content updates: editorial changes go through the database with no code release needed

RevupAds Dashboard

React 19 • PHP • MySQL • Recharts

- Multi-tenant ad-revenue analytics platform: admin uploads CSV or XLSX reports, clients see per-website dual-axis Recharts
- Session-cookie auth with per-website SQL scoping ensuring clients only access their own data