

CONTACT

Paris, France
+33 7 75 73 96 00
chmudasir.fr@gmail.com
LinkedIn • GitHub
Portfolio

TECHNICAL COMPETENCIES

Programming
Python, SQL
TypeScript, JavaScript

Data Engineering
Airflow, Kafka, Spark, Hadoop
dbt, ETL/ELT, streaming
Data Vault 2.0

Databases and Warehouses
PostgreSQL, BigQuery
Snowflake, Redshift, Databricks
Redis, pgvector

Cloud
AWS, GCP, Azure
S3, Lambda, Step Functions
Vercel, Cloudflare R2

Governance and Quality
Lineage, validation, versioning
audit history, unit normalization
reference mappings

Backend and APIs
FastAPI, REST, Django, SQLAlchemy
Alembic, Pydantic, JWT, RBAC
NextAuth, Celery, Socket.IO

Frontend
Next.js, React, Tailwind CSS

DevOps and Infrastructure
Docker, Kubernetes, Terraform
GitHub Actions, CI/CD, Nginx
Linux, SSH, Let's Encrypt

AI and Automation
OpenAI, Claude, Gemini
PydanticAI, scikit-learn, n8n

Analytics and Observability
Power BI, Tableau, Grafana, GA4
PostHog, Sentry, Prometheus

Sustainability Data
LCA, carbon accounting
product/BOM and process models
factor mapping, lineage, versioning

EDUCATION

MSc Data Engineering
Aivancity University Paris | 2023-2024
BS Mathematics
Air University | 2016-2020

CERTIFICATIONS

AWS Cloud Practitioner • IBM Data Engineer
Google Data Analytics • Machine Learning - Coursera

LANGUAGES

English - Advanced | French - A2/B1

Mudasir Shafiq

LEAD DATA ENGINEER | CLOUD DATA PLATFORMS | GOVERNANCE AND AI

Data Engineer with 4+ years of experience owning the architecture and delivery of cloud data platforms, batch and streaming pipelines, analytical models, and AI-enabled data products. Strong in Python and SQL across Airflow, Kafka, dbt, Spark, PostgreSQL, BigQuery, Redshift, Snowflake, AWS, GCP, and Azure. Builds secure, traceable systems with data-quality controls, lineage, versioned reference data, RBAC, CI/CD, and production monitoring.

EXPERIENCE

Data and AI Engineer
Jul 2026-Present
Growlytics | Remote

- Own end-to-end data architecture for a multi-tenant competitive-intelligence platform, including PostgreSQL modelling and serverless ingestion across Meta Ad Library, Google Ads Transparency, TikTok via Apify, Meta Marketing, and Google Ads APIs.
- Built deterministic performance-metric pipelines and grounded Claude/Gemini analysis, separating computed evidence from generated narratives for traceable outputs.
- Automated scheduled ingestion, report generation, billing, and email delivery with Vercel, Blob storage, cron workflows, Paddle, and Resend.

Data and AI Engineer / Full Stack Engineer
Jun 2026-Present
GrowMore and Collabs+ | Dubai / Remote

- Designed data-intensive services with Next.js, FastAPI, PostgreSQL, SQLAlchemy, Alembic, Pydantic, JWT/RBAC, Redis caching, file ingestion, and background email workflows.
- Modelled Collabs+ across 9 business modules and 35 PostgreSQL tables, enforcing constraints, role-based access, soft deletion, immutable audit history, and notification preferences.
- Defined scalable AI matching and real-time flows with pgvector/IVFFlat, Celery/Redis, Socket.IO Pub/Sub, Cloudflare R2, OpenAI, and PydanticAI.
- Established containerized delivery, isolated environments, CI/CD, health monitoring, and secure Nginx/TLS deployment architecture.

Data Engineering Consultant
Feb 2026-Jul 2026
Watershed | Contract | Paris / Remote

- Designed normalized models connecting products, bills of materials, manufacturing processes, and environmental factors for traceable product-impact calculations.
- Built repeatable ETL workflows that standardized units, validated keys and mappings, and preserved raw-to-normalized lineage across spreadsheet and catalogue data.
- Implemented versioned reference mappings and quality checks for ambiguous candidates, missing factors, unit mismatches, and reproducible downstream calculations.

Data Engineer
Dec 2023-Jun 2024
Ceres-Advisory - Blockchain and Web3 | Paris

- Architected a real-time GA4-to-BigQuery pipeline through GTM, reducing reporting latency from hours to minutes across 15+ KPI dashboards.
- Optimized multi-million-row blockchain data with partitioning and window functions, reducing SQL query time by 60%.
- Designed resilient AWS ingestion with Lambda, Step Functions, and S3, achieving zero post-deployment pipeline failures.
- Developed a DistilBERT sentiment pipeline and surfaced model outputs through Power BI and Grafana monitoring views.

Data Analyst and Performance Marketing
Jan 2022-Sep 2023
PriceOye - E-commerce | Pakistan

- Sustained ROAS of 6.33x-43x and delivered PKR 18M monthly GMV through data-driven campaign optimization and audience analysis.
- Raised retention from 10% to 15% through AI-powered audience segmentation and automated lifecycle campaigns.
- Automated Power BI pipelines and KPI dashboards, eliminating more than 10 hours of manual reporting each week.

SELECTED DATA ENGINEERING PROJECTS

Livestock Intelligence | ***Kafka • Airflow • BigQuery • scikit-learn***

- Built a Kafka-to-BigQuery pipeline with Airflow orchestration, validation, CTE feature engineering, and LAG window functions.
- Prepared model-ready features for health classification, fertility prediction, and milk-yield regression.

Reddit ETL Pipeline | ***Airflow • Terraform • AWS • dbt • Redshift • Tableau***

- Automated AWS ingestion with Airflow and Terraform, then modelled Data Vault 2.0 transformations in dbt and Redshift.
- Published curated analytical datasets through repeatable Tableau dashboards for stakeholder reporting.

Synthetic Data Engine | ***Python • Validation Rules • Google Sheets***

- Generated configurable daily datasets with field-level constraints and business-rule validation.
- Implemented idempotent Google Sheets upserts to prevent duplicates and support repeatable downstream testing.